

# Self-guided Approximate Linear Programs

---

Last updated on: [May 10, 2024](#)

**Topics:** [approximate dynamic programming](#), [approximate linear programming](#), [model-based reinforcement learning](#), [random Fourier features](#), [inventory management](#), [options pricing](#).

---

## Introduction

---

This repository hosts algorithm implementations and benchmarks discussed in the paper **Self-Guided Approximate Linear Programs: Randomized Multi-Shot Approximation of Discounted Cost Markov Decision Processes** authored by [Parshan Pakiman](#), [Selva Nadarajah](#), [Negar Soheili](#), and [Qihang Lin](#), available at [Management Science](#). Specifically, this repository includes implementations of the following algorithms for computing control policies in Markov decision processes (MDPs):

- Feature-based Approximate Linear Program (FALP)
- Self-guided FALP
- Policy-guided FALP
- Least Squares Monte Carlo ([Longstaff and Schwartz 2001](#))

Furthermore, the repository includes implementations of two algorithms for computing lower bounds on the optimal policy cost of MDPs:

- Information Relaxation and Duality ([Brown et al. 2010](#))
- A heuristic based on Constraint Violation Learning ([Lin et al. 2020](#))

Integrating lower bounds with the upper bounds derived from simulating control policies enables the computation of optimality gaps. The optimality gap of a policy is the difference between the cost of this policy and the lower bound, expressed as a percentage of the lower bound. This repository implements MDPs associated with the following applications:

- Perishable inventory control
- Bermudan options pricing

**Note:** The code in this repository is relatively general and can be extended to other MDPs to compute control policies and lower bounds.

## Instruction

---

The following steps are tailored for macOS and Ubuntu users. Windows users may adapt and create similar instructions suitable for their operating system.

1. Download the [repository](#) on your local system and extract the zip file. For ease of reference, we assume the extracted files are stored in the following path: /Desktop/Self-Guided-ALPs-Discounted-Cost-master.

2. Open Terminal on your machine and run the following code:

```
cd /Desktop/Self-Guided-ALPs-Discounted-Cost-master
```

3. Please check the version of your Python. For example, run the code below:

```
python3 --version
```

4. Please confirm that Python 3.10 or above is installed on your machine. There are different ways to install Python. We leave this step to the user's discretion.
5. Create a virtual Python environment and name it ALP. For example, use the following code:

```
python -m venv ALP
```

6. Activate the ALP environment as follows:

```
source ALP/bin/activate
```

7. This code relies on several Python packages. We utilize pip for their installation, but users can explore alternative methods. If choosing pip, please use the following code for its update:

```
pip install --upgrade pip
```

8. Please install the following libraries on the ALP environment:

- numpy
- pandas
- scipy
- numba
- tqdm
- emcee
- sampl
- importlib
- sampl\_mcmc
- nengo
- gurobipy

For instance, run the following command to install all the above libraries:

```
python -m pip install numpy pandas scipy numba tqdm e
```

9. This repository uses [Gurobi](#) for solving large-scale linear programs.

Please ensure that Gurobi (`gurobipy`) is installed along with its corresponding license. If you are affiliated with academia, Gurobi provides a free academic license, accessible through [this page](#). In some cases, an error might occur when running the `grbgetkey` code to activate your license. Refer to [this](#) troubleshooting page for assistance with resolving this issue. If the issue persists, consider installing Gurobi using `conda` instead of `pip`. You can use the following code as an example:

```
conda install -c gurobi gurobi  
grbgetkey xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx
```

10. Provided that the ALP environment and Gurobi are correctly installed, you can use the following code to solve test instances of perishable inventory control applications. Please run the following code:

```
./run_PIC.sh
```

The `run_PIC.sh` file solves the first instance of the perishable inventory control problem using the FALP algorithm with 20 random Fourier features, where FALP is formulated using a uniform state-relevance distribution (refer to the paper and code for details). You can find the specifications (e.g., cost parameters) of this instance under

the path MDP/PIC/Instances/instance\_1.py. To solve this instance using an alternate algorithm, you can modify the file run\_PIC.sh. For instance, changing the algo\_name from FALP to SG-FALP in this file and rerunning run\_PIC.sh will display the output for the self-guided FALP algorithm applied to this instance. A screenshot of the output of these algorithms is attached below:

```

=====
Instance number: 1
Algorithm name: FALP
Basis function: Fourier
State relevance: uniform_non_adaptive
# inner updates: 5
Random basis seed: 111
Optimality threshold: 0.0%
=====
# Basis | # Constrs | FALP Obj | ALP ConT | ALP SltV | Lower Bound | LB RT | Policy Cost | UB RT | Opt Gap (N) | TOT RT |
-----
20 | 200000 | 2122.2 | 0.104 | 0.086 | 2883.4 | 0.399 | 2881.6 | 0.543 | 2.0 | 1.284 |
=====
Instance number: 1
Algorithm name: SGFALP
Basis function: Fourier
State relevance: uniform_non_adaptive
# inner updates: 5
Random basis seed: 111
Optimality threshold: 0.0%
=====
# Basis | # Constrs | FALP Obj | SGFALP Obj | ALP ConT | ALP SltV | SG T | Lower Bound | LB RT | Policy Cost | UB RT | Opt Gap (N) | TOT RT |
-----
5 | 200000 | 1796.7 | 1796.7 | 0.188 | 0.000 | 0.000 | 1711.2 | 0.229 | 2495.7 | 0.574 | 31.0 | 0.911 |
10 | 200000 | 1948.6 | 1948.6 | 0.188 | 0.032 | 0.000 | 1872.6 | 0.318 | 2564.8 | 0.564 | 57.0 | 1.975 |
15 | 200000 | 2121.2 | 2121.2 | 0.186 | 0.037 | 0.078 | 1994.0 | 0.375 | 2879.6 | 0.572 | 4.0 | 3.136 |
20 | 200000 | 2121.2 | 2121.2 | 0.111 | 0.078 | 0.137 | 2082.9 | 0.467 | 2864.6 | 0.578 | 2.0 | 4.908 |
=====

```

11. Running file run\_BerOpt.sh produces the output of the algorithms for Instance 1 of the Bermudan options pricing.

## Replication

The code in this repository can be directly used to replicate the tables and figures in the paper. For perishable inventory control problem instances, we consider four methods:

- $ALP^{LNS}$ : the standard ALP model formulated using basis functions described in [Lin et al. 2020](#).
- FALP: a randomized ALP model described in §3 of the paper.
- $FALP^{PG}$ : policy-guided FALP model described in Algorithm 1 of the paper.
- $FALP^{SG}$ : self-guided FALP model described in Algorithm 2 of the paper.

For the Bermudan options pricing problem, we consider methods FALP and  $FALP^{SG}$ , in addition to the following ones:

- LSM: Least squares Monte Carlo ([Longstaff and Schwartz 2001](#)) that is implemented in file LSM.py.
- $ALP^{DFM}$ : the standard ALP model formulated using basis functions described in [Desai et al. 2012](#).

We test these methods on different sets of instances described in the paper. To obtain the tables and figures in the paper, the parameters of two files, run\_PIC.sh and run\_BerOpt.sh, should be modified. We compute the lower and upper bounds for each instance-method pair across ten different

trials. Each of these trials is specified by the value of parameter **seed** that is varied in {111, 222, 333, 444, 555, 666, 777, 888, 999, 1010}. For each method-instance-trial, we compute the lower and upper bounds and then the averages of these bounds across the trials for each method-instance pair. Below, we describe the other parameters in `run_PIC.sh` and `run_BerOpt.sh` to obtain each table and figure in the paper.

- **Table 2:** Compares  $ALP^{LNS}$  and  $FALP$  with on 12 instances. To replicate this table, please run `run_PIC.sh` after modifying the following parameters:

Parameter name in <code>run_PIC.sh</code>	$ALP^{LNS}$	$FALP$
<code>algo_name</code>	FALP	FALP
<code>basis_func_type</code>	Ins	fourier
<code>max_basis_num</code>	150	150
<code>batch_size</code>	150	150
<code>state_relevance_inner_itr</code>	0	0
<code>instance_number</code>	{1, 2, ..., 12}	{1, 2, ..., 12}

Please note that `main.py` automatically sets the value of parameters **max\_basis\_num** and **batch\_size** for each instance.

- **Table 3:** Compares four models  $ALP^{LNS}$ ,  $FALP$ ,  $FALP^{PG}$ , and  $FALP^{SG}$  on 6 instances. To replicate this table, please run `run_PIC.sh` after modifying the following parameters:

Parameter name in <code>run_PIC.sh</code>	$ALP^{LNS}$	$FALP$	$FALP^{PG}$	$FALP^{SG}$
<code>algo_name</code>	FALP	FALP	PG-FALP	SG-FALP
<code>basis_func_type</code>	Ins	fourier	fourier	fourier
<code>max_basis_num</code>	300	300	300	300
<code>batch_size</code>	300	300	50	300
<code>state_relevance_inner_itr</code>	0	0	0	5
<code>instance_number</code>	{13, 14, ..., 18}	{13, 14, ..., 18}	{13, 14, ..., 18}	{13, 14, ..., 18}

- **Table 4:** Compares four models  $ALP^{LNS}$ ,  $FALP$  with 600 random basis functions,  $FALP$  with 1000 random basis functions, and  $FALP^{SG}$  on 6 instances. To replicate this table, please run `run_PIC.sh` after modifying the following parameters:

Parameter name in <code>run_PIC.sh</code>	$ALP^{LNS}$	$FALP_{600}$	$FALP_{1000}$	$FALP^{SG}$
<code>algo_name</code>	FALP	FALP	PG-FALP	SG-FALP
<code>basis_func_type</code>	lms	fourier	fourier	fourier
<code>max_basis_num</code>	600	600	1000	600
<code>batch_size</code>	600	600	1000	100
<code>state_relevance_inner_itr</code>	0	0	0	0
<code>instance_number</code>	{19, 20, ..., 24}	{19, 20, ..., 24}	{19, 20, ..., 24}	{19, 20, ..., 24}

- **Figure 2:** Depicts the violin plot of upper and lower bounds across 10 trials for 6 iterations of algorithm  $FALP^{SG}$  on two instances: Instance 19 and Instance 20. Upon running this algorithm on these instances, several files will be generated under the paths `Output/PIC/instance_19` and `Output/PIC/instance_20`. Then, one can use the plotting function in `SG_FALP_progress_plot.py` to obtain Figure 2 based on these files.
- **Table 5:** Presents performance of LSM,  $ALP^{DFM}$ ,  $FALP$ , and  $FALP^{SG}$  on 9 instances. To obtain this table, please modify file `run_BerOpt.sh` by letting **instance\_number** take values in {1, 2, ..., 9} and **instance\_number** take values in {1, 2, ..., 9} and **seed** to take values in {111, 222, 333, 444, 555, 666, 777, 888, 999, 1010}.

## Appendix

This code has undergone testing on two systems:

- 2021 MacBook Pro:
  - CPU: M1 Pro
  - Memory: 16 GB

- OS: Mac OS 14.1.2
- 2022 Mac Studio:
  - CPU: M1 Max
  - Memory: 64 GB
  - OS: Mac OS 14.1.1

Below is the list of Python package versions utilized during the testing phase of this repository:

Package	Version
autograd	1.6.2
emcee	3.1.4
future	0.18.3
gurobipy	11.0.0
importlib	1.0.4
llvmlite	0.41.1
nengo	4.0.0
numba	0.58.1
numpy	1.26.2
pandas	2.1.4
pip	23.3.1
python-dateutil	2.8.2
pytz	2023.3.post1
sampyl-mcmc	0.3
scipy	1.11.4
setuptools	63.2.0
six	1.16.0
tqdm	4.66.1
tzdata	2023.3